

Declarative Integration of Interactive 3D Graphics into the World-Wide Web: Principles, Current Approaches, and Research Agenda

Abstract

With the advent of WebGL, plugin-free hardware-accelerated interactive 3D graphics has finally arrived in all major Web browsers. WebGL is an imperative solution that is tied to the functionality of rasterization APIs. Consequently, its usage requires a deeper understanding of the rasterization pipeline. In contrast to this stands a declarative approach with an abstract description of the 3D scene. We strongly believe that such approach is more suitable for the integration of 3D into HTML5 and related Web technologies, as those concepts are well-known by millions of Web developers and therefore crucial for the fast adoption of 3D on the Web. Hence, in this paper we explore the options for new declarative ways of incorporating 3D graphics directly into HTML to enable its use on any Web page. We present declarative 3D principles that guide the work of the *Declarative 3D for the Web Architecture W3C Community Group* and describe the current state of the fundamentals to this initiative. Finally, we draw an agenda for the next development stages of Declarative 3D for the Web.

CR Categories: I.3.7 [Computer Graphics]: Three-Dimensional Graphics and Realism—Virtual Reality I.3.6 [Methodology and Techniques]: Standards—Languages

Keywords: Declarative 3D, HTML5, DOM Integration, Polyfill

1 Introduction

The Web evolved from a text-based system to the current rich and interactive medium that supports images, 2D graphics, audio and video. These types of new media have made the Web experience richer, more attractive to users, etc., than ever before, and opened up possibilities for new types of applications and usage. The major media type that is still missing is 3D: synthetic, possibly photorealistic images in 3D with animation, as smoothly integrated in the everyday Web experience as images or video. Just as the appearance of images or video could open new application possibilities, access to the 3D on a Web site would make it possible to include realistic models of 3D objects – from models of buildings to representation of the human body or the sceneries for computer games. With WebGL [Marrin 2012], a JavaScript binding for OpenGL ES 2.0, this seems feasible; however, the goal would be to achieve the same smooth inclusion of 3D content in a Web page like we experience today with images or SVG-based 2D graphics.

Although some of these goals could also be achieved by imperative means (e.g., through the usage of WebGL), developments of 3D models have a long tradition of using declarative approaches, which is also in line with some of the fundamental principles of Web development. It is therefore important to explore how the experiences accumulated in two different communities, namely the Web Development and Computer Graphics communities, can be capitalized upon to achieve the long term goal of using 3D on the Web the same way as we do with 2D graphics and video today. Moreover, while imperative graphics APIs are powerful and necessary, a *Declarative 3D* approach can provide web authors an easy way to add interactive high-level declarative 3D objects through the

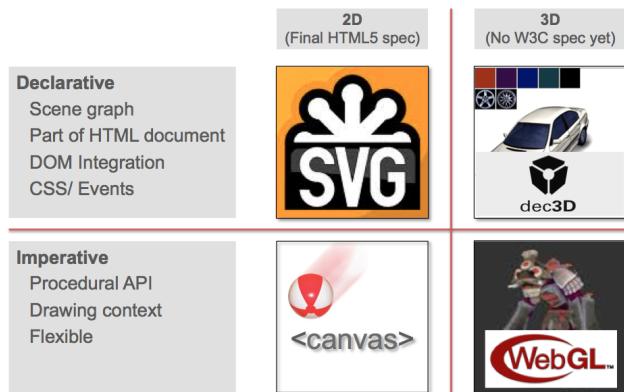


Figure 1: The position of the Declarative 3D approach in the current Web graphics technology ecosystem.

HTML Document Object Model (DOM) allowing them to easily create, modify, share, and experience interactive 3D graphics using HTML documents. Figure 1 depicts the position of the Declarative 3D approach in the current Web graphics technology ecosystem.

It is arguable that the emerging support for an imperative 3D API for the Web is useful but insufficient for broad acceptance and usage of 3D on the Web. A declarative approach that offers qualified concepts and that is tightly integrated with current web technologies, such as JSON or XML for scene construction, DOM for scene manipulation, and CSS for styling, is necessary to support a fast adoption and broad use of interactive 3D graphics by the millions of Web developers. The provided concepts must lift the hardware-oriented imperative application programming interfaces to an expressive and more easily usable level. Therefore not the low-level data structures of existing GPU layers must be in the center of the design but high-level elements and items like 3D objects, transformations, material descriptions, and lights. Instead of teaching Web developers 3D graphics APIs, the goal is to bring 3D graphics to the point where it is natural for Web developers to just make use of it. While this might not be possible for every possible use of a low-level API, it still can cover the vast majority of use cases.

The *Declarative 3D for the Web Architecture W3C Community Group* (Dec3D) [W3C Community Group 2013] was thus founded to suggest and create methods to add high-level declarative 3D objects to the HTML DOM [W3C 2005], so users can share and experience interactive 3D contents. Moreover, this not only allows creating new content from existing content but also to index and search 3D content. The core mission is to determine the requirements, options, and use cases for the integration of interactive 3D graphics into the Web technology stack in a declarative way, which hopefully will provide a foundation for future standardization. Therefore, the group aims at presenting common use cases that define how 3D might intersect and interact with HTML5, DOM events, CSS, SVG, GeoLocation, Augmented Reality, Efficient XML Interchange (EXI) [W3C 2011] and other key working groups, whereas certain complex data types (e.g., transformation matrices) and computations are also of mutual interest. In this regard, this paper presents the current state and efforts concerning Dec3D.

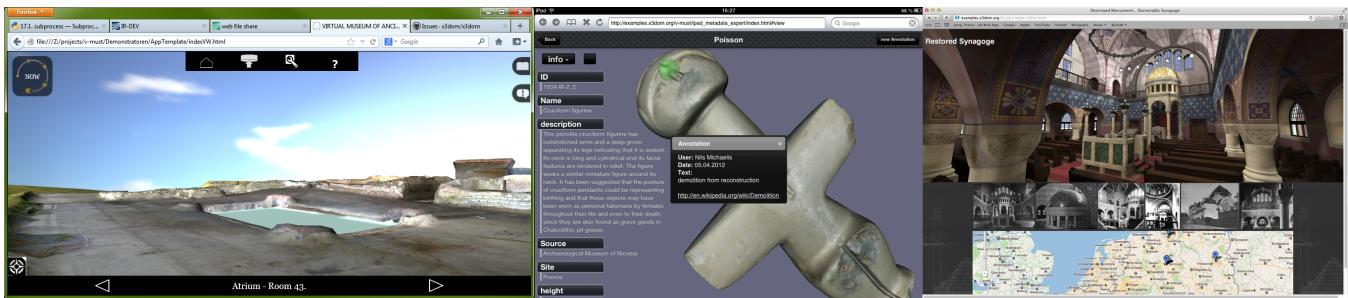


Figure 2: From left to right: walkthrough of Roman heritage site; cultural heritage object explorer with metadata; restored virtual synagogue.

2 Declarative 3D Principles

Here we describe declarative 3D principles, where the following goals should guide the development of DOM-based 3D graphics.

Following the Established Principles of the Web Declarative 3D is being developed to significantly lower the barrier for authoring 3D content for Web sites by duplicating the key features that enabled the growth of the early Web and its further success.

Separation of structure from content Underlying the Web from its earliest days was the separation of structure from content. The concept of a paragraph specified by the `<p>` tag was separate from its content. Declarative 3D is attempting to bring the same separation to 3D graphics inside of web pages. The concepts such as definitions of 3D objects, transformations, materials, etc. should be implemented in a declarative 3D description as an extension to HTML5 [Hickson 2012] using any existing or future extension mechanism.

Separation of content from style One of the principles of the current Web is also the separation of content and style, most notably through CSS. The successful integration of SVG [W3C 2012b] with HTML was made much easier due to the fact that SVG was already following this principle. The objective here is to extend the use of CSS for styling 3D graphics. One example can be the use of the latest CSS3 3D Transforms [W3C 2012a] to allow manipulating not only 2D but also 3D objects.

Use of the Document Object Model The DOM [W3C 2005] is a platform- and language-neutral interface that allows programs and scripts to dynamically access and update the content, structure and style of Web documents. Declarative 3D should use the DOM API to examine and modify elements on the 3D scene and their attributes by simply reading and setting their properties. As the DOM provides access to user actions (e.g., pressing a key or clicking a mouse button), it should also be used as a main interface to interact with 3D contents.

Embedded 3D graphics should reuse existing W3C techniques, specifically from HTML5 and SVG, as far as possible and propose extensions only where specific features are necessary or provide significant benefits. Where new concepts are introduced their relation to and effects on existing Web standards should be analyzed, evaluated, and discussed with the respective W3C working groups.

3D Content Creation and Reuse While the creation of original 3D geometry and appearances still requires 3D specific know-how, the reuse, configuration, and manipulation of such content should be made similarly easy as for 2D Web content now. The solution should hide internal data structures and algorithms and provide users convenient ways to edit and manipulate such scenes. A key success factor for Declarative 3D on the Web will be the ability to

generate new or reuse existing content. This requires that suitable exporters and converters can be built. However, as 3D on the Web is supposed mainly as a delivery mechanism, it is not necessary to include the ability to semantically represent all 3D features.

Platform Independence 3D content needs to be described in a way that does rely neither on a specific render API such as OpenGL or DirectX nor on a specific rendering technique such as rasterization or ray tracing only [Schwenk et al. 2012]. This should allow for content to be portable across browsers, rendering techniques, and hardware platforms, while taking advantage of available features wherever possible. The results of rendering content under such different environments should be highly predictable.

Efficiency and Scalability Interactive real-time 3D graphics enables new forms of interactivity on the Web but also adds significant new requirements on user agents. A key requirement for the selected technology therefore is the possibility to implement it efficiently (cf. e.g. [Trevett 2012]). Native implementations allow utilizing all (battery) resources more efficiently while leveraging heterogeneous hardware. Thereby, the CPU time can be used for the application instead for rendering, collision, scene-housekeeping, etc., which is esp. critical for mobile devices. Since 3D scenes can become rather large, any solution should target scalability in the sense that 3D content should run across different platforms (from mobile devices to high-end graphics hardware) with predictable performance. Mechanisms should be in place to handle cases where the performance provided by a user agent on some platform is not sufficient, e.g. by allowing for switching to different content (e.g. lower LOD) or provide alternate methods of delivering the content (e.g. server-based rendering delivered via streaming video).

Security and Digital Rights Management Secure delivery of Web content is a general problem and not specific to 3D data. However, the economic value of 3D data might make the problem more acute. Any proposed solution should therefore be based on a general approach to secure Web content. However, we first need to collect use cases, extract requirements and examine how far existing methods (e.g. [Koller and Levoy 2005]) and standards can be transferred to the proposed architecture. It is already demonstrated that the application of XML Encryption and Signature is needed for document fragments as well as full documents, since high-fidelity or sensitive portions of 3D models often need special protections.

Accessibility and Usability Accessibility improvements serve all users, not just people with disabilities. A problematic aspect of many 3D graphics approaches however is that user navigation and interaction is implemented inconsistently. Therefore, users familiar with one approach are impeded when navigating or interacting with other 3D scenes and models. Examination of relevant Web Accessi-

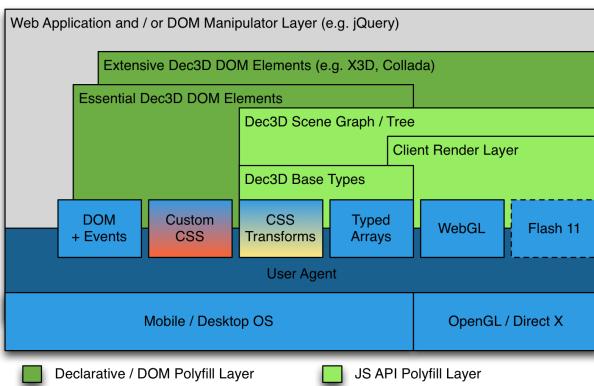


Figure 3: Proposed declarative 3D “polyfill” runtime architecture.

bility Initiatives (WAI) principles might provide significant benefit. Conversely, use of declarative 3D graphics models might provide major benefits when describing the accessibility features and constraints of real-world objects and locations. Declarative 3D goals and potential solutions may achieve significant benefits if they are harmonizable with WAI imperatives.

Leveraging Web Development Infrastructure The Web has become the new application development platform for our time. Countless rich internet applications are being created at a dizzying pace. The “apps” world for mobile applications has also become a huge phenomenon. One important enabling technologies is JavaScript [Crockford 2008], which has evolved from a toy language into a robust and richly nuanced tool that is the basis for much of the recent explosion in Web applications. Declarative 3D, by living within the structure of Web apps is poised to leverage the huge collection of tools and infrastructure that already exists such as jQuery [Resig 2012] etc. The ability to take existing Web debuggers, editors, and viewers and use these tools to create, edit, and debug 3D content is a tremendous benefit and just begins to scratch the surface of development tools, created for Web application development, but which we can use for 3D content development.

3 Declarative 3D Frameworks

As already mentioned, the Dec3D Community Group [W3C Community Group 2013] has been formed to determine the requirements, options, and use cases for the declarative integration of interactive 3D graphics capabilities into the Web technology stack, which will provide a foundation for future standardization. While this standardization is our goal, we still need platforms allowing for the experimentation and for the assessment of our design decisions. We also need to reach out to Web developers who could provide us with valuable feedback as early as possible.

3.1 X3DOM and XML3D

With X3DOM [Behr et al. 2009] and XML3D [Sons et al. 2010] two open-source frameworks are available, which are briefly introduced in this section. Both frameworks support the ongoing discussion in the computer graphics and Web communities how an integration of HTML(5) and Declarative 3D content could look like. This discussion is followed by the presentation of the so-called “polyfill” approach for further development activities.

Fraunhofer IGD’s X3DOM [Behr et al. 2009; Behr et al. 2012b]

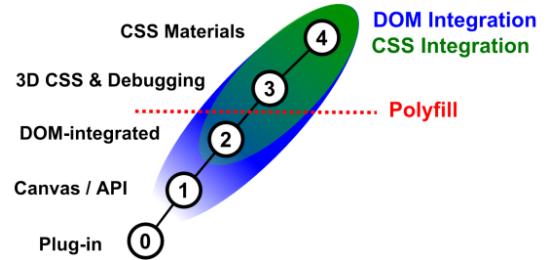


Figure 5: Levels of integration: current browser APIs only allow integrating with DOM and DOM Events, but not yet fully with CSS.

is a JavaScript-based open-source framework for declarative 3D graphics in HTML5 that aims at extending the HTML DOM tree with declarative 3D objects while employing modern Web technologies like CSS3, Ajax, DOM scripting, as well as WebGL and – as fallback – Adobe’s Flash 11 with Stage 3D [Adobe 2013] for GPU-accelerated rendering. The proposed 3D elements are mostly based upon the open ISO standard X3D [Web3D Consortium 2011], though X3DOM introduces a special HTML profile that basically extends the X3D Interchange profile. Additionally, instead of implementing the somehow odd X3D pointing device sensor component, X3DOM simply uses, and appropriately extends, the HTML UI/Mouse events such that 3D pick events are likewise supported.

Furthermore, to overcome various problems that come along when embedding 3D mesh data, which typically consists of several megabytes of vertex attribute data, directly into the DOM tree, the developers are working on efficient 3D mesh encodings that allow separating the node structure from the raw vertex data [Behr et al. 2012b]. In this regard, Figures 2 and 4 show several X3DOM applications from the cultural heritage and engineering domain, where the ability to efficiently handle big 3D data sets is of high importance. Based on this work, 3D transmission formats are now also a major topic at the Khronos Group [Trevett 2012], where it is similarly argued that raw 3D data, just like image, video or audio data, needs to be externalized from the HTML document.

DFKI’s XML3D [Sons et al. 2010] is a similar approach. In contrast to X3DOM, this format proposes a new tag set that is not based on X3D anymore. Furthermore, it does not define a scene-graph but a strict tree structure by means of a prototyping mechanism. In recent versions of XML3D, mesh data can also be externalized from the HTML document by storing it in separate JSON files. Moreover, the developers are working on a system for dynamic data-flow to efficiently support animations [Klein et al. 2012].

3.2 Polyfill Approach

We also propose to measure a level of integration for 3D graphics in terms of W3C technologies (DOM, CSS, etc.) to explain existing and possible integration levels (Figure 5 shows that currently only level 2, DOM integration, is possible), and to finally identify a suitable integration level for declarative 3D. However, the latter imposes several problems in that it requires having a set of 3D-specific CSS properties or even a full shader description via CSS as well as a deeper integration into debugging tools like Chrome Developer Tools or Firebug. However, integration levels 3 and 4 are currently not yet achievable by a JavaScript implementation, because the required APIs are missing: it is not possible to introduce new CSS properties or to extend developer tools by an API.

It quickly became clear that a higher level of integration requires additional APIs in user agents. After discussions with browser vendors (namely Firefox and Chrome), who made it clear that integra-

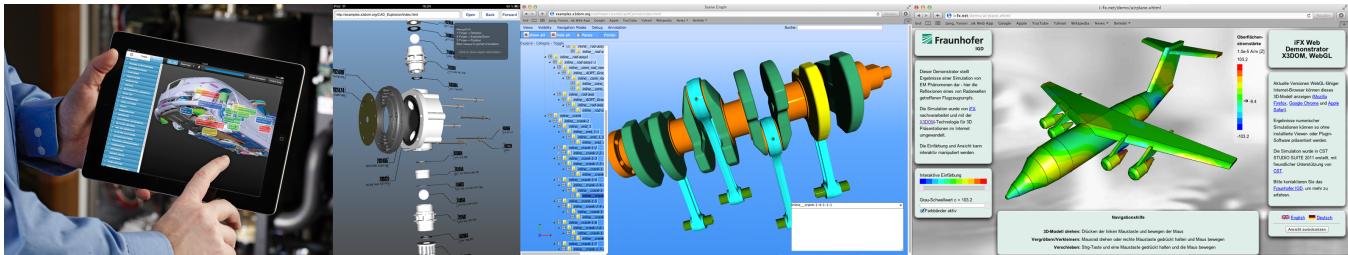


Figure 4: From left to right: web-based design review application with annotation markers; interactive explosion of CAD model on iPad; lightweight web-based viewer for 3D CAD data; visualization and interactive exploration of simulation data by using extended mouse events.

tion of the desired extensions natively into their frameworks is not currently their priority, we decided to consider a polyfill approach for the further development activities. Nevertheless, communicating and requesting additional and missing features like custom CSS properties or CSS monitoring is necessary.

A *polyfill* is a downloadable piece of code, which provides facilities that are not natively built-in to a Web browser [Sharp 2010]. For example, many features of HTML5 are not supported by versions of Internet Explorer older than version 8 or 9, but can be used by web pages if those pages install a so-called polyfill. Polyfills can also be used to add entirely new functionality to browsers.

Polyfill-based approaches allow deriving hard requirements for related and utilized W3C standards and user agent APIs immediately. This leads to much more evolving concepts and solutions in contrast to an overall declarative 3D specification. To sum up, the X3DOM and XML3D experimental declarative 3D Web publishing frameworks are designed to explore different options for adding 3D graphics to HTML. Here, Figure 3 depicts our proposed declarative 3D polyfill runtime architecture, though we'd like to stress that the whole integration model is still evolving and open for discussion.

3.3 Declarative 3D Essentials

In this section we briefly outline our proposed essential elements for an upcoming Dec3D standard. First of all, a start element allows the transition from the HTML flow layout to 3D transformations:

```
294 <div>
295   <dec3d style="border: 1px solid black;">
296   ...
297   </dec3d>
298 </div>
```

The DOM integration obviously extends the HTML hierarchy. But this comes along with certain issues like the correct treatment of CSS inheritance. Therefore, CSS integration is another point in that existing CSS properties should be used where applicable. Transformations for instance can be represented using CSS 3D Transforms [W3C 2012a]. In addition, a set of 3D-specific CSS properties (e.g. for defining the appearance) needs to be defined. However, custom CSS properties are not yet supported by any Web Browser, though finally this seems to change now [W3C 2013].

Moreover, not only standard DOM UI Events like ‘mouseover’ and ‘click’ shall be supported, but also events that contain 3D-specific context information. In this regard, the Browser API should expose convenience functions and special Dec3D base types consisting of a lean set of complex data types for 2D/3D graphics. Similarly, generic mesh data containers should be based on TypedArrays [Khronos 2012] with appropriate interfaces to efficiently modify the data, where external binary containers are used as transmission format [Behr et al. 2012b; Trevett 2012]. Thus, 3D geometry definition will stay close to OpenGL primitives.

Besides this, we further distinguish between essential Dec3D DOM Elements, were the focus lies on a minimalistic element set, as well as more comprehensive Dec3D DOM Elements (see middle of Figure 3), which build on top of those with the focus on usability and which for convenience could build on existing scene-graph standards like X3D [Web3D Consortium 2011].

Essential DOM elements are a basic set of built-in shaders, which should be enough for most use cases. Advantages are that they do not entail timing attack issues and that they can adapt to any target device. As textures, HTML elements like ``, `<video>`, `<canvas>` shall be used, whereas elements like `<svg>` and `<html>` can serve as interactive textures.

Lights are basically also shader parameters in OpenGL/WebGL [Marrin 2012], but here shall be represented as own elements to provide intuitive ways to define lights and to not obstruct more advanced future concepts. Likewise, viewpoints can serve as link to certain points of interest within the document. However, it is still discussable if some pre-defined camera navigation modes should be part of the essential profile.

4 Declarative 3D Agenda

In this section we shape the agenda and identify upcoming research issues for the next development stage of Dec3D. During the 1st Intl. Workshop on Declarative 3D [Behr et al. 2012a] as well as during more informal meetings at Web3D and SIGGRAPH in 2011/2012, the members and supporters of the Declarative 3D W3C Community Group agreed upon the following topics to make the effort successful and for the W3C to adopt/develop a Dec3D standard.

Encourage Participation All relevant stakeholders, e.g. developers, designers, researchers, 3D artists, industry professionals, representatives of standards organizations, accessibility experts, and user-agent implementers, are encouraged to participate in this group. Participants should be willing to actively develop and donate materials towards the group's deliverables, as well as attend the group's teleconferences and face-to-face meetings.

Clear Definition of Use Cases and Requirements The group needs to agree on a collection of use cases, where embedding 3D data in HTML using a declarative approach provides significant benefit. Here, declarative approaches esp. seem relevant for the industry, since they tend to think in formats not in APIs. Each use case should explore how publishers and consumers benefit from Dec3D. From these use cases, the group needs to derive and prioritize the different required dimensions for the Dec3D technical specification [Jankowski 2012; Le Feuvre 2012].

Clear Technical Specification The next step is the creation of the clear, detailed and extensible technical specification of the im-

363 plementation concepts and features necessary to cover a majority
 364 of useful requirements. Measurable properties need to be defined
 365 to quantitatively and/or qualitatively evaluate the achieved solution,
 366 document the pros and cons of each, and demonstrate that, based on
 367 the above analysis, there is a good chance of success in creating a
 368 W3C standard for Declarative 3D for the Web.

369 **Outreach and Exemplar Applications** The group needs to con-
 370 tinue its outreach activities through the high quality demonstra-
 371 tions of the Declarative 3D philosophy using the open-source
 372 frameworks X3DOM and XML3D. Therefore, several applications
 373 should be identified, each requiring and demonstrating different ca-
 374 pabilities of Dec3D. For example, one application could require
 375 huge 3D datasets which are impractical for inclusion directly into
 376 the DOM; another could require real-time control of an external
 377 system (inter-process communication); another could integrate with
 378 a complex data base; another an illustration of complex data (infor-
 379 mation visualization) without a physical analog, and another a more
 380 traditional scientific visualization. Key to selection of these applica-
 381 tions is the use and demonstration of a declarative 3D requirement.

382 **W3C Working Group Proposal** Finally, the Community Group
 383 should deliver reports documenting its progress, any conclusions
 384 it arrived at with respect to the standardization of *Declarative 3D*
 385 for the Web and, if reaching a positive conclusion, recommending
 386 a standardization approach as a basis for a future W3C working
 387 group on the same topic.

388 5 Conclusions

389 While WebGL, a 3D imperative graphics API in the Web context,
 390 is getting more and more traction, we are still missing an easy way
 391 to add interactive high-level declarative 3D objects to an HTML
 392 document to allow anyone to easily create, share, and experience
 393 interactive 3D graphics, with possibly wide ranging effects similar
 394 to those caused by the broad availability of video on the Web. The
 395 main motivation is thus to make it easy to add 3D to Web pages
 396 by bringing 3D to the Web developers and not vice versa. This
 397 can be achieved by fully integrating 3D content into HTML5 docu-
 398 ments, where interactive 3D graphics is a first class DOM objects.
 399 Moreover, by reusing existing Web technology wherever possible,
 400 no new concepts are added unless absolutely necessary.

401 Another objective of this position paper on Declarative 3D for the
 402 Web is to evaluate the options for a successful standardization of
 403 a declarative approach to interactive 3D graphics as part of HTML
 404 documents. The idea is to collect suitable use cases, derive require-
 405 ments from them, and then find the essential set of features and con-
 406 cepts that enables broad uptake by authors and users of interactive
 407 3D on the Web. We are absolutely aware that our goal is ambitious
 408 and it will take some time to implement these features. Therefore,
 409 we call for more participation from the Web3D and W3C commu-
 410 nities that we believe is crucial to achieve our common and ultimate
 411 goal: 3D for everyone and everywhere.

412 References

- 413 ADOBE, 2013. Adobe flash player 11. <http://www.adobe.com/products/flashplayer.html>. 414
- 415 BEHR, J., ESCHLER, P., JUNG, Y., AND ZÖLLNER, M. 2009. 416 X3DOM – a DOM-based HTML5/ X3D integration model. In
 417 *Proc. Web3D 2009*, ACM Press, New York, USA, 127–135.
- 418 BEHR, J., BRUTZMAN, D., HERMAN, I., JANKOWSKI, J., AND 419 SONS, K., Eds. 2012. Proceedings of the 1st Intl. Work-
 420 shop on Declarative 3D for the Web Architecture (Dec3D2012
 421 at WWW2012), vol. 869 of *CEUR Workshop Proceedings*.
- 422 BEHR, J., JUNG, Y., FRANKE, T., AND STURM, T. 2012. Using
 423 images and explicit binary container for efficient and incremental
 424 delivery of declarative 3d scenes on the web. In *Proc. Web3D*,
 425 ACM, New York, NY, USA, 17–25.
- 426 CROCKFORD, D. 2008. *JavaScript: The Good Parts*. O'Reilly.
- 427 HICKSON, I., 2012. HTML5 (W3C Working Draft).
<http://www.w3.org/TR/2012/WD-html5-20120329/>.
- 428 JANKOWSKI, J. 2012. Writing effective use cases for the declarative 3d for the web architecture. In *Dec3D'12*, CEUR-WS.
- 429 KHRONOS, 2012. Typed array spec. <http://www.khronos.org/registry/typedarray/specs/latest/>.
- 430 KLEIN, F., SONS, K., JOHN, S., RUBINSTEIN, D., SLUSALLEK, P., AND BYELOZYOROV, S. 2012. Xflow: declarative data processing for the web. In *Web3D*, ACM Press, 37–45.
- 431 KOLLER, D., AND LEVOY, M. 2005. Protecting 3d graphics content. *Commun. ACM* 48, 6, 74–80.
- 432 LE FEUVRE, J. 2012. Towards declarative 3d in web architecture. In *Dec3D'12*, CEUR-WS.
- 433 MARRIN, C., 2012. WebGL specification. <https://www.khronos.org/registry/webgl/specs/latest/>.
- 434 RESIG, J., 2012. jQuery. <http://jquery.com/>.
- 435 SCHWENK, K., JUNG, Y., VOSS, G., STURM, T., AND BEHR, J. 2012. CommonSurfaceShader revisited: Improvements and experiences. In *Proceedings Web3D 2012: 17th Intl. Conf. on 3D Web Technology*, ACM Press, New York, USA, 93–96.
- 436 SHARP, R., 2010. What is a polyfill? <http://remysharp.com/2010/10/08/what-is-a-polyfill/>.
- 437 SONS, K., KLEIN, F., RUBINSTEIN, D., BYELOZYOROV, S., AND SLUSALLEK, P. 2010. Xml3d: interactive 3d graphics for the web. In *Web3D '10*, ACM.
- 438 TREVETT, N. 2012. 3d transmission format. In *Seventh AR Standards Community Meeting Talks*.
- 439 W3C COMMUNITY GROUP, 2013. Declarative 3D for the Web Architecture. <http://www.w3.org/community/declarative3d/>.
- 440 W3C, 2005. Document Object Model (DOM). <http://www.w3.org/DOM/>.
- 441 W3C, 2011. Efficient xml interchange (exi) format. <http://www.w3.org/TR/2011/REC-exi-20110310/>.
- 442 W3C, 2012. Css 3d transforms. <http://dev.w3.org/csswg/css3-3d-transforms/>.
- 443 W3C, 2012. Svg. <http://www.w3.org/Graphics/SVG/>.
- 444 W3C, 2013. Css custom properties for cascading variables module level 1. <http://dev.w3.org/csswg/css-variables/>.
- 445 WEB3D CONSORTIUM, 2011. Extensible 3d (X3D). <http://www.web3d.org/x3d/specifications/>.